

## A detour into the world of Linux schedulers

### **Abstract:**

I explored 3 different schedulers implemented in Linux and how they perform on different processes. We tested Round Robin, First-In-First-Out and the CFS scheduler with I/O bound processes, CPU-bound processes and a mixture of those processes and charted their performance. I found that CFS faltered when it came to IO-intensive processes whilst outperforming RR and FIFO on CPU-intensive processes. RR and FIFO performed relatively similarly.

### **Introduction:**

In this report, I explored various schedulers available within the Linux OS; round robin (SCHED\_RR), FIFO (SCHED\_FIFO) and CFS (SCHED\_OTHER). A scheduler is crucial to the running of a modern OS. Schedulers are often implemented so they keep all compute resources busy (as in load balancing), allow multiple users to share system resources effectively, or to achieve a target quality of service. Scheduling is fundamental to computation itself, and an intrinsic part of the execution model of a computer system; the concept of scheduling makes it possible to have computer multitasking with a single central processing unit (CPU).<sup>1</sup>

The tests were performed on a 64-bit Ubuntu 12.04 VM allocated 5728 MB with 4 cores of a 2.7 GHz Intel Core i5.

### **Method:**

To thoroughly test the schedulers, we needed to establish certain processes. I created a CPU-intensive process that calculates pi and an I/O-bound process that reads an input file and writes data to a number of outputs. I then combined these into a testing C file that requires several parameters to perform its task. We pass in the scheduler policy (S, the number of processes we will fork and the type of processes we'd like to test. Given these parameters, our program runs the required task.

e.g.

```
$ ./test-sched SCHED_FIFO SMALL IO
$ ./test-sched SCHED_OTHER MEDIUM CPU
```

```
<SCHEDULER>: SCHED_FIFO, SCHED_RR, SCHED_OTHER (DEFAULT)
<# OF ITERATIONS>: LOW (DEFAULT) (10), MEDIUM(100), HIGH(1000)
<PROCESSTYPE>: CPU, IO, MIXED
```

To test our programs and their combinations systematically, we wrote a bash script that iterates through all of the combinations presented. We had 3 schedulers, 3 amounts of processes and 3 types of process which results in  $3^3$  combinations, or 27 unique combinations. We ran the tests 5 times each to account for discrepancies and average them out. We run 500,000 iterations of every process. The bash script tests them in order and returns data in a

---

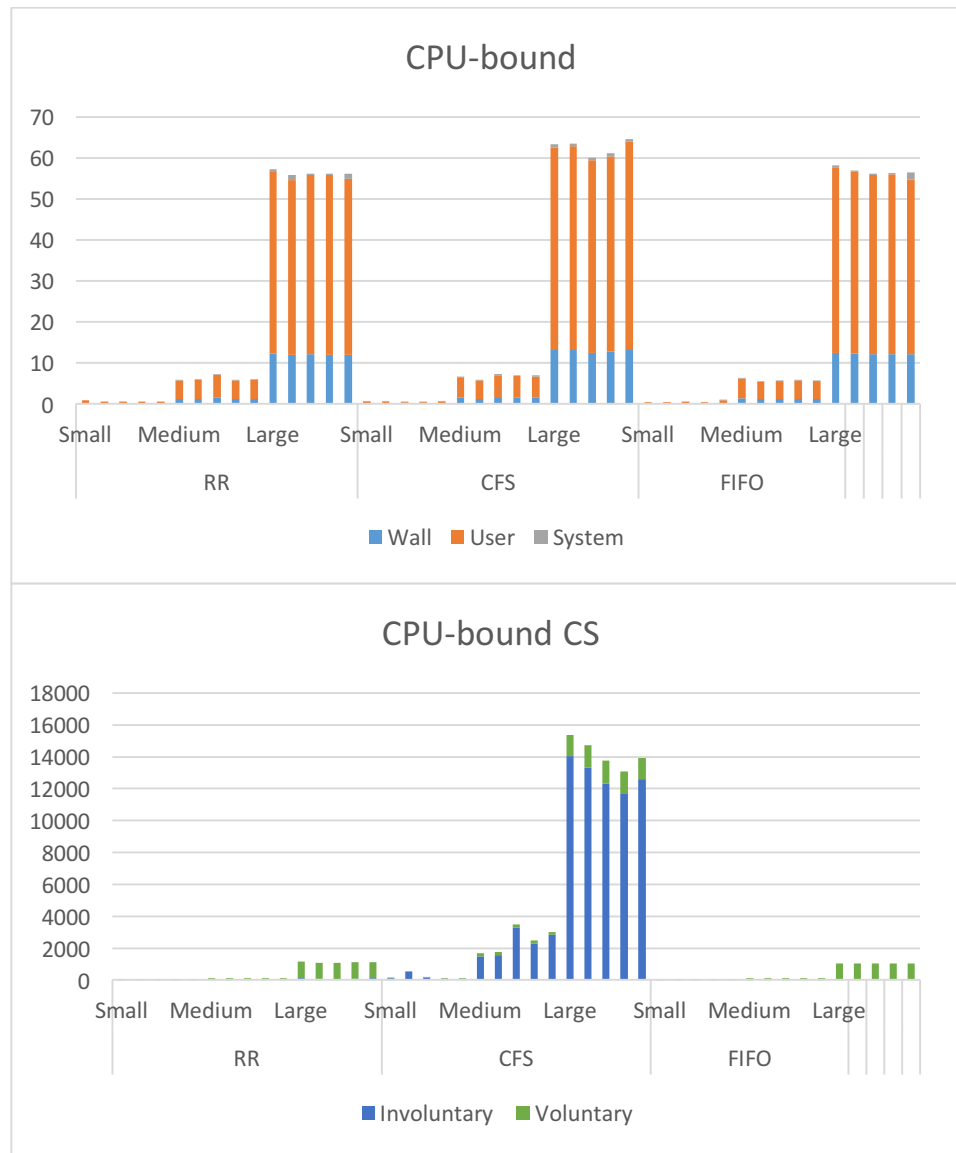
<sup>1</sup> [https://en.wikipedia.org/wiki/Scheduling\\_\(computing\)](https://en.wikipedia.org/wiki/Scheduling_(computing))

text file for easy visualization using spreadsheet software. Without further ado, let's move onto the results section!


### **Results:**

The graphs are vertical GANTT charts. We show the run times of the processes in a chart and the amounts of voluntary/involuntary context switches in the following chart.

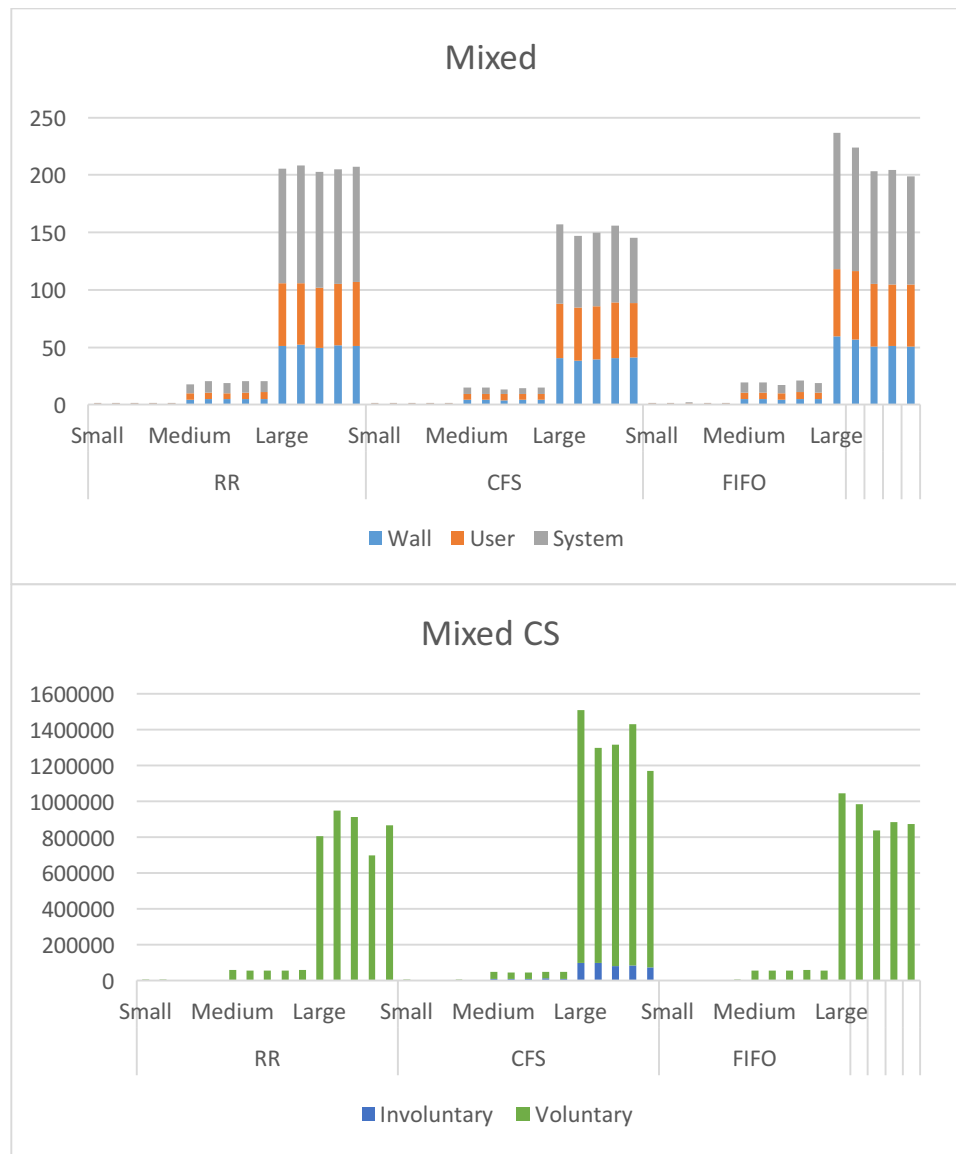
CPU-intensive:



IO-Bound:



Mixed:



### Analysis:

To begin, we'll focus on the CPU-bound processes. We notice that all three of the schedulers performed almost equally well. However, looking at the context switch chart for CPU-bound processes shows a radically different picture. The CFS spends relatively massive amounts doing involuntary context switches. They're understandably involuntary, as the CPU's won't yield to the scheduler whilst there's still a task to be completed. Despite those large wait times, it still performed on almost equal footing as RR and FIFO (fractionally slower). We conclude that RR and FIFO are marginally better at CPU-intensive tasks.

On the other hand, in the I/O bound processes we notice that despite having more than double the amount of context switches, the CFS is about 40% quicker than both RR and FIFO which again are almost equal. The majority of the context switches are voluntary, as expected

when dealing with I/O devices. In this case, CFS outperforms both RR and FIFO scheduling and is our clear winner.

Lastly, our mixed process results are predictable given our previous results. Since CFS is almost on par with RR and FIFO in CPU-Intensive processes and 40% quicker in I/O bound processes, then our mixed process results show that CFS is about 25% quicker than both RR and FIFO despite have 25-30% more context switches.

### **Conclusion:**

In conclusion, we see that CFS outperforms both RR and FIFO schedulers in I/O bound processes and mixed processes which most accurately reflect personal real-world applications. Both RR and FIFO are on par with each other and outperform CFS in CPU-intensive processes, though suffer serious drawbacks with I/O bound processes.

Building off our results, CFS would be most efficient for modern-day personal computing whilst RR/FIFO would work best in scientific computing that utilize heavy CPU usage with little I/O sequences. RR would falter heavily in interactive applications, which CFS can handle. CFS would increase wait times substantially given a really CPU intensive scientific computation that spans weeks/years.

We notice that CFS doesn't scale very well in the CPU-intensive tasks, with the wait times increasing near exponentially while RR/FIFO understandably remain linear in their growth.

### **References:**

Andy Sayler  
Class Lectures  
Operating Systems Concepts – Silverschatz

### **Appendix A:**

	Scheduler	Processes	Wall	User	System	CPU	Involuntary	Voluntary
CPU	RR	Small	0.21	0.65	0	309%	3	18
			0.14	0.46	0	321%	4	18
			0.15	0.48	0	320%	4	18
			0.14	0.46	0	331%	3	18
			0.15	0.49	0	324%	5	18
		Medium	1.22	4.5	0.01	368%	6	107
			1.26	4.63	0.02	368%	9	108
			1.54	5.56	0.04	363%	12	108
			1.25	4.57	0.02	367%	9	108
			1.26	4.68	0.02	371%	8	109
		Large	12.29	44.54	0.38	365%	127	1041
			12.04	42.64	1.19	364%	62	1009
			12.13	43.79	0.32	363%	58	1011
			12.09	43.74	0.35	364%	88	1025

	CFS	Small	12.08	42.93	1.13	364%	134	1008
			0.18	0.5	0.01	277%	135	24
			0.18	0.39	0.01	218%	520	23
			0.14	0.36	0.01	251%	164	24
			0.12	0.38	0.01	322%	64	23
			0.13	0.43	0.01	342%	74	24
		Medium	1.53	5.06	0.12	337%	1493	193
			1.3	4.53	0.08	354%	1566	197
			1.77	5.32	0.23	313%	3303	201
			1.58	5.29	0.18	347%	2305	199
			1.58	5.19	0.24	342%	2834	187
		Large	13.46	49.11	0.84	371%	14032	1330
			13.38	49.49	0.64	374%	13337	1382
			12.56	46.9	0.56	377%	12334	1450
			12.77	47.59	0.77	378%	11712	1389
			13.49	50.43	0.72	379%	12598	1314
	FIFO	Small	0.12	0.39	0	330%	4	18
			0.12	0.39	0	316%	6	18
			0.14	0.46	0	316%	5	18
			0.12	0.4	0	323%	6	16
			0.22	0.72	0.02	331%	10	17
		Medium	1.34	4.88	0.04	366%	6	108
			1.18	4.29	0.02	365%	8	108
			1.2	4.4	0.04	368%	9	108
			1.22	4.49	0.03	369%	8	108
			1.23	4.41	0.02	358%	10	107
		Large	12.47	45.26	0.41	366%	48	1008
			12.28	44.4	0.3	364%	53	1008
			12.1	43.8	0.36	364%	48	1009
			12.12	43.85	0.34	364%	52	1008
			12.13	42.62	1.67	364%	50	1008
IO	Scheduler	Processes	Wall	User	System	CPU	Involuntary	Voluntary
	RR	Small	0.41	0.01	0.61	150%	6	4015
			0.51	0	1.07	208%	6	3401
			0.44	0	0.8	183%	5	3772
			0.34	0	0.72	209%	4	3811
			0.36	0	0.65	179%	5	3622
		Medium	3.35	0.06	8.76	262%	15	55247
			3.8	0.01	10.38	273%	19	60543

			3.47	0.05	9.36	271%	20	59868
			3.79	0.02	10.3	272%	29	55824
			3.66	0.02	9.29	254%	19	56583
		Large	38.65	0.34	101.24	262%	502	623868
			35.68	0.26	97.56	274%	469	697109
			39.31	0.32	106.16	270%	429	620903
			37.56	0.28	102.97	274%	415	708244
			37.04	0.3	102.08	276%	416	679669
	CFS	Small	0.32	0	0.31	99%	561	2360
			0.28	0	0.3	105%	529	3092
			0.2	0	0.18	90%	291	2991
			0.37	0.01	0.38	107%	397	2791
			0.19	0	0.1	55%	258	2566
		Medium	2.52	0.02	3.64	145%	7636	37344
			2.88	0.02	4.5	157%	8363	36996
			2.58	0	3.97	154%	7599	40435
			2.47	0	3.65	147%	7712	38795
			2.84	0.02	4.64	164%	8252	39979
		Large	28.05	0.1	58.98	210%	63009	1297221
			27.59	0.12	64.68	234%	73939	1557439
			27.99	0.12	67.13	240%	63745	1628215
			27.2	0.18	57.28	211%	79247	1191711
			28.75	0.2	66.61	232%	76711	1323476
	FIFO	Small	0.37	0	0.66	177%	5	3639
			0.37	0	0.62	165%	6	3640
			0.59	0	1.26	213%	3	3498
			0.48	0	1	209%	5	3723
			0.34	0	0.64	190%	7	4207
		Medium	3.52	0	9.42	267%	18	55853
			3.72	0.05	9.41	253%	11	56884
			3.38	0.05	8.9	265%	11	56899
			3.54	0.06	10.1	287%	13	55935
			3.17	0.05	8.34	264%	11	64608
		Large	40.84	0.27	109.29	268%	94	648636
			41.58	0.31	109.6	264%	130	651701
			38.08	0.32	99.02	260%	109	603797
			38.64	0.23	102.92	266%	102	609987
			36.96	0.28	97.68	265%	116	605897
Mixed	Scheduler	Processes	Wall	User	System	CPU	Involuntary	Voluntary

	RR	Small	0.43	0.56	0.3	198%	3	4837
			0.37	0.56	0.22	208%	8	4477
			0.51	0.54	0.5	205%	4	3578
			0.39	0.56	0.31	219%	5	3927
			0.37	0.49	0.26	203%	6	4172
		Medium	4.5	5.58	7.61	292%	65	60114
			5.03	5.38	10.08	307%	64	56209
			4.6	5.37	8.98	311%	49	55996
			4.91	5.75	9.67	313%	58	57532
			4.85	5.91	9.51	318%	73	57920
		Large	51.33	54.66	99.82	300%	651	803791
			52.24	53.74	102.53	299%	684	946067
			49.71	52.32	100.6	307%	687	911065
			51.69	53.78	99.79	297%	630	697404
			51.15	55.65	100.46	305%	684	866909
	CFS	Small	0.48	0.51	0.33	176%	852	2972
			0.49	0.5	0.46	192%	684	2930
			0.4	0.46	0.31	194%	556	3113
			0.4	0.46	0.46	228%	1036	3108
			0.54	0.56	0.69	232%	715	2923
		Medium	4.16	5.04	5.47	252%	9792	39081
			4.31	5.11	5.36	242%	11540	33295
			3.71	5.35	4.15	256%	9051	37001
			4.12	5.18	5.31	254%	11599	37835
			4.27	4.94	5.76	250%	10717	37622
		Large	40.49	47.38	69.45	288%	100158	1409292
			38.43	46.15	62.58	282%	99389	1198817
			39.24	46.38	64.13	281%	81040	1235369
			40.77	48.28	67.09	282%	83025	1348386
			40.86	47.36	57.02	255%	73862	1095514
	FIFO	Small	0.48	0.54	0.64	247%	6	3926
			0.46	0.56	0.6	247%	6	3853
			0.53	0.52	0.88	265%	4	4032
			0.38	0.47	0.5	255%	5	3866
			0.36	0.52	0.28	217%	5	4656
		Medium	4.83	5.78	8.82	302%	11	55200
			4.84	5.54	9.04	301%	16	56786
			4.36	5.34	7.59	296%	11	55481
			5.1	5.68	10.15	310%	15	60247



		4.61	5.66	8.34	303%	10	54628
	Large	59.74	58.1	118.67	295%	139	1042543
		56.47	59.81	107.4	296%	149	984459
		50.66	54.39	98.38	301%	156	837963
		51.13	53.29	100.15	300%	145	881949
		50.4	54.02	94.31	294%	161	873904

**Appendix B:**

- test-sched.c :- runs the various processes required for testing
- testscript :- runs a systematic analysis of all 27 cases